

Introduction to Scikit-learn

Objective:

This assignment will help you get hands-on experience with Scikit-learn's basic functionalities, including loading datasets, data preprocessing, splitting datasets, training models, evaluating models, and tuning hyperparameters.

```
In [24]: import warnings
import sys
if not sys.warnoptions:
    warnings.simplefilter("ignore")
```

Step 1: Load and Explore the Dataset

```
In [26]: from sklearn.datasets import load_iris
import pandas as pd

# Load the Iris dataset
data = load_iris()

# Convert to a Pandas DataFrame
df = pd.DataFrame(data.data, columns=data.feature_names)
df['target'] = data.target

# Display dataset details
print("Shape of dataset:", df.shape)
print("Feature names:", data.feature_names)
print("Target classes:", data.target_names)

# Display the first 5 rows
print(df.head())
```

Shape of dataset: (150, 5)

Feature names: ['sepal length (cm)', 'sepal width (cm)', 'petal length (cm)', 'petal width (cm)']

Target classes: ['setosa' 'versicolor' 'virginica']

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	\
0	5.1	3.5	1.4	0.2	
1	4.9	3.0	1.4	0.2	
2	4.7	3.2	1.3	0.2	
3	4.6	3.1	1.5	0.2	
4	5.0	3.6	1.4	0.2	

	target
0	0
1	0
2	0
3	0
4	0

Step 2: Preprocess the Dataset

In [28]: *# Use StandardScaler to scale the features so they have a mean of 0 and a standard deviation of 1.*

```
from sklearn.preprocessing import StandardScaler

# Features and Target
X = df.iloc[:, :-1] # Features
y = df['target']     # Target

# Scale the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

print("Scaled features (first 5 rows):\n", X_scaled[:5])
```

Scaled features (first 5 rows):

```
[[-0.90068117  1.01900435 -1.34022653 -1.3154443 ]
 [-1.14301691 -0.13197948 -1.34022653 -1.3154443 ]
 [-1.38535265  0.32841405 -1.39706395 -1.3154443 ]
 [-1.50652052  0.09821729 -1.2833891  -1.3154443 ]
 [-1.02184904  1.24920112 -1.34022653 -1.3154443 ]]
```

Step 3: Split the Dataset

```
In [30]: from sklearn.model_selection import train_test_split

# Split the dataset
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)

print("Training set size:", X_train.shape[0])
print("Testing set size:", X_test.shape[0])
```

Training set size: 120

Testing set size: 30

Step 4: Train a Random Forest Classifier

```
In [32]: from sklearn.ensemble import RandomForestClassifier

# Initialize and train the model
model = RandomForestClassifier(random_state=42)
model.fit(X_train, y_train)

print("Random Forest model trained successfully.")
```

Random Forest model trained successfully.

Step 5: Evaluate the Model

```
In [34]: from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, classification_report

# Predict on test data
y_pred = model.predict(X_test)

# Calculate evaluation metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred, average='weighted')
recall = recall_score(y_test, y_pred, average='weighted')
f1 = f1_score(y_test, y_pred, average='weighted')

print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1-Score:", f1)
```

```
# Detailed classification report
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

Accuracy: 1.0
 Precision: 1.0
 Recall: 1.0
 F1-Score: 1.0

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	9
2	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

Step 6: Tune Hyperparameters Using Grid Search

In [36]: `from sklearn.model_selection import GridSearchCV`

```
# Define hyperparameter grid
param_grid = {
    'n_estimators': [50, 100, 150],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10]
}

# Initialize GridSearchCV
grid_search = GridSearchCV(RandomForestClassifier(random_state=42), param_grid, cv=5, scoring='accuracy')
grid_search.fit(X_train, y_train)

# Best parameters and score
print("Best Parameters:", grid_search.best_params_)
print("Best Accuracy:", grid_search.best_score_)

# Evaluate the tuned model on test data
best_model = grid_search.best_estimator_
```

```
y_pred_tuned = best_model.predict(X_test)
print("Accuracy after tuning:", accuracy_score(y_test, y_pred_tuned))
```

Best Parameters: {'max_depth': None, 'min_samples_split': 5, 'n_estimators': 50}

Best Accuracy: 0.95

Accuracy after tuning: 1.0

Step 7: Visualize Feature Importance

```
In [38]: import matplotlib.pyplot as plt
import numpy as np

# Get feature importances
feature_importances = model.feature_importances_

# Plot feature importances
plt.barh(np.array(data.feature_names), feature_importances)
plt.xlabel('Importance')
plt.ylabel('Features')
plt.title('Feature Importance in Random Forest')
plt.show()
```

